


Benchmark context, without mixing the evidence


Separate populations, checkpoints and planner RNG. These bars are descriptive, not cross-study treatment effects.

A AUTHOR-REPORTED Reach


CEM L2 · Tables 2/11; ‡final checkpoint from Table 12




Reach-Wall




Push-T




PointMaze



Wall




DROID† score




B HISTORICAL DEVELOPMENT

Earlier six-task measurements · not the protected panel


Reach




Reach-Wall




Push-T




PointMaze



Wall




DROID† score




C FRESH PROTECTED

96 new scenarios/task × all 8 arms · concurrent native within each scenario


Reach




Reach-Wall



PointMaze



Wall



* Historical MetaWorld Joint / Visual / Action used native 45.83 / 29.17, not the canonical 44.79 / 30.21 bars.

† DROID is recorded-action agreement, not physical success. Point estimates only; no comparable error bars. All 48 fresh paired intervals include zero.