
Context-Consistent Steering for Planning in JEPA World Models

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 An internal edit can change a world model’s forecast without changing the selected
2 action, and a patch that matches an input counterfactual at one step can diverge as
3 context advances. We study both effects in frozen JEPA-WM checkpoints using
4 low-rank predictor corrections, action-conditioning patches, and controlled plan-
5 ning replays. Across sixteen development contexts, replacing an action-derived
6 condition at both appearances in a two-frame context reproduces the raw-action
7 counterfactual exactly through six forecast steps when all predictor blocks are
8 patched. Replacing only its first appearance loses roughly half of the reconstruc-
9 tion by the final step. Separately, a learned rank-four correction improves pro-
10 prioceptive forecast MSE by 2.36% and 2.19% on two tasks. A score-margin
11 bound certifies unchanged winners in 184 of 192 fixed candidate banks, while
12 adaptive search can still return different plans. Physical replay and a 3,072-run
13 protected evaluation do not establish improved control performance. These re-
14 sults distinguish counterfactual consistency, forecast accuracy, and planning util-
15 ity. The context study tests repeated action history in a short rolling window; it
16 does not measure long-token context or context-length scaling.

17 1 Introduction

18 A predictive world model supports planning by forecasting the consequences of candidate actions.
19 Improving its predictions is useful to a planner when the change helps it select actions with better
20 outcomes. To study this connection, we keep JEPA-WM checkpoints frozen and intervene inside
21 their action-conditioned predictor [3]. Low-rank corrections, visual and action-conditioning edits,
22 and matched random controls let us measure changes in predicted futures, goal scores, selected
23 plans, and physical execution separately.

24 Sequential prediction adds another question: does an internal intervention keep representing the in-
25 tended change as the model consumes its evolving context? In the predictor studied here, an action
26 enters two successive context windows. Replacing its internal representation once need not repro-
27 duce changing the input action itself. We compare one-time and history-consistent donor patches
28 with a coherent raw-action substitution, holding model weights and candidate banks fixed. Patching
29 both appearances at every block reproduces the specified counterfactual through H6; patching only
30 the first appearance matches initially, then diverges.

31 This is a concrete issue for interventions in agents that repeatedly forecast, act, and update their
32 context. Our empirical scope is a two-frame predictor with six-step imagined rollouts and receding-
33 horizon planning over longer simulator episodes. We do not enlarge the model’s context window
34 or test long-token retrieval. The connection to longer-horizon agents is the need to check how an
35 intervention behaves when the edited information is reused, rather than assuming its first-step effect
36 persists.

How JEPA-WM plans with predicted futures

Frozen JEPA-WM · alternative intervention sites at imagined step H3




Figure 1: Steering inside a JEPA world model. The predictor forecasts the consequences of candidate action sequences. The planner scores their distance to the encoded goal, refits its action distribution, and repeats. It executes the beginning of the chosen plan before observing and replanning. R marks the four-direction predictor edit; V and A mark the alternative visual and action-conditioning edits. All edits occur at imagined step H3 of an H6 forecast. R is never combined with V/A in the eight-condition evaluation.

Our representational questions are informed by Joseph et al.’s Physics Emergence Zone analysis [1], which combines layerwise probes, subspace analyses, and targeted interventions in video models. Here the target is an action-conditioned world model whose predictions drive a sampling planner. The contributions are a controlled action-history test, a fixed-bank bound connecting score perturbations to plan selection, and an evaluation that follows forecast corrections through adaptive search and physical execution. Layer sweeps and numerical controls qualify the interpretation of the internal edits. These experiments use separate populations; they do not establish one causal chain explaining protected behavior.

2 Related work

DINO-WM and JEPA-WM use prediction in learned features for planning [3, 5]. JEPA-WM studies which model and training choices support physical planning. We keep the released models fixed and ask what an inference-time activation correction changes.

Activation steering also appears in robot policies. COAST uses success- and failure-derived concepts in a policy’s action expert [2]. Our intervention enters a world-model predictor whose output is scored by a separate sampling planner. The supervision and computations differ, so comparing reported gains does not isolate the reason for their difference.

Joseph et al. probe physical variables and intervene on motion representations in video models [1]. Their layer analyses motivate checking where an internal signal is accessible and what changes when it is edited. Our six-layer action-conditioned predictor is a different object of study, and its attention maps do not identify a physics-emergence zone. Curved activation interpolation [4] motivates a second control: reconstruction quality should be tested against downstream prediction and numerical precision before receiving a physical interpretation.

59 3 Model and experimental setup

60 **Prediction and planning.** Frozen encoders map images and robot state to feature vectors. A six-
 61 block predictor takes these features and candidate actions to forecast six imagined steps. In the
 62 MetaWorld tasks, the planner minimizes

$$C(a_{1:H}) = \text{MSE}(\hat{z}_H^v, z_g^v) + 0.1 \text{MSE}(\hat{z}_H^p, z_g^p), \quad (1)$$

63 where v and p denote visual and robot-state features, and each MSE averages over its feature dimen-
 64 sions. The target is the goal, whereas forecast error compares a prediction with the outcome of those
 65 actions.

66 The Cross-Entropy Method (CEM) starts with 300 action sequences, keeps the ten with lowest pre-
 67 dicted cost, and samples a new population around them. MetaWorld uses fifteen such iterations
 68 per planning call. The robot executes a three-step prefix, corresponding to fifteen elementary sim-
 69 ulator actions, before replanning. Full MetaWorld episodes contain seven planning calls over 100
 70 elementary steps; navigation uses thirty search iterations per call.

71 **Interventions.** The main learned edit adds a normalized correction in a four-dimensional subspace
 72 to the fourth predictor block’s output (B3), at the third imagined step (H3). Its coefficients come
 73 from an offline fitted map of the current unedited features. A random-subspace comparison matches
 74 rank, edit site, and prescribed dose, with separate response calibration. Both are activation edits; the
 75 robot still obtains its actions from the planner.

76 We also test a learned pair of visual-input and action-conditioning directions. Equal-budget and
 77 randomized versions form one comparison. Unscaled joint, visual-only, and action-only edits form
 78 a component ablation. Together with unsteered and the two four-direction conditions, these give
 79 eight arms. Appendix A gives the equations, fitting details, and dose definitions. No intervention
 80 updates the base model weights.

81 **Data and uncertainty.** Fitting trajectories determine the operators. Separate recorded develop-
 82 ment trajectories measure prediction error and inform the recipe. Development diagnostics compare
 83 fixed plans, iterative searches, and short physical executions. The protected evaluation uses 96 new
 84 scenarios per task on Reach, Reach-Wall, PointMaze, and Wall after freezing the eight-arm recipe.
 85 All arms share the same physical GPU within each scenario, with strict FP32 and TF32 disabled.
 86 Engineering checks use excluded inputs.

87 Statistical units are source trajectories for recorded forecasts and starting scenarios for planning
 88 and behavior. Candidate plans are not independent replicates. Registered uncertainty uses paired
 89 bootstrap resampling and correction across the specified comparisons: six for the search extension,
 90 twelve for physical replay, and 48 for protected behavior. Each uses 20,000 draws. Details and
 91 source identities are retained in the appendices and accompanying reports. Push-T and DROID
 92 appear only in earlier development results; DROID measures recorded-action agreement.

93 4 Results

94 4.1 An action patch must follow the action through context

95 We replace the H3 action with a donor action and use the resulting raw-input rollout as the counter-
 96 factual reference. The predictor’s action encoder maps 20 action coordinates into a 400-dimensional
 97 condition. At H3 the changed action occupies the newest context position; at H4 it occupies the
 98 older position. We compare condition replacement at H3 only with replacement at both H3 and
 99 H4. The experiment retains sixteen development contexts (eight each on Reach and Reach-Wall),
 100 original and freshly seeded candidate banks, and all six predictor blocks, for 1,120 H6 forecasts.
 101 Candidate banks share contexts and are not independent samples.

102 Let $\hat{z}_{\text{cf}}$ denote the coherent raw-action forecast and $\hat{z}_{\text{native}}$ the unmodified forecast. At each mea-
 103 sured endpoint we report

$$R = 1 - \frac{D(\hat{z}_{\text{patch}}, \hat{z}_{\text{cf}})}{D(\hat{z}_{\text{native}}, \hat{z}_{\text{cf}})}, \quad (2)$$

An action change has to survive its own history.

Replace it at H3, then replace its historical occurrence at H4.




Figure 2: Matching a changed action now does not ensure matching its later forecast. Upper panels retain all-block patches; lower panels show B1 as an explicitly post hoc illustration of a single-block edit. Both tasks, both banks, and all sixteen contexts are retained. The dotted line is the coherent raw-action reference. Bands are existing marginal 95% context-bootstrap intervals, not simultaneous layer comparisons. The all-block persistent patch matches exactly; the B1 persistent patch does not. Only H3, H4, and H6 are measured here.

where D averages candidate-wise visual MSE plus 0.1 times proprioceptive MSE before taking the ratio. $R = 1$ is exact counterfactual reconstruction and $R = 0$ is the unmodified forecast's error. This measures agreement with another model rollout, not with the physical future.

Replacing both appearances at all six blocks reproduces the raw-action substitution byte-exactly through H6. Replacing all blocks only at H3 matches at H3, then diverges when the original action returns as history at H4. Its H6 reconstruction is $R = 0.5029$ – 0.5079 across the four task/bank cells (Figure 2). Exact all-block parity is a computational consistency control: replacing every occurrence reconstructs the changed input's computation. The one-time patch measures what is lost when that replacement is incomplete.

Single-block interventions give a depth-dependent response. Persistent minus H3-only reconstruction is positively separated at B0–B3 in both tasks and both banks: sixteen of twenty-four cells under the registered simultaneous correction. B4/B5 remain unresolved. B1 has the largest observed persistence difference, $\Delta R = 0.4991$ – 0.5057 ; its persistent H6 reconstruction is 0.7320 – 0.7421 , versus 0.2286 – 0.2387 for H3-only. B1 is an illustrative post hoc choice, not a newly selected steering site. Appendix C.4 retains all six layers, input-range controls, and all 72 primary contrasts.

The mismatch reaches candidate selection. A post hoc reanalysis of the saved H6 costs compares each patch with the coherent raw-action reference. For all-block H3-only patches, mean Spearman agreement is 0.878 – 0.894 and mean top-ten overlap is 5.75 – 6.88 across task/bank cells. Their winning candidate differs from the reference in 17 of 32 banks (two banks per context); persistent all-block patches agree exactly. B1 persistence improves mean rank agreement and elite overlap in every task/bank cell, while winner agreement improves in three of four. These descriptive comparisons retain all sites and contexts. They measure fixed-bank selection, not a new CEM or physical rollout (Appendix C.5).

For world-model intervention studies, the practical consequence is to validate an intended input counterfactual across the rolling context, not just at the first edited step. The experiment does not show that every useful activation edit must correspond to a raw-action change. Nor does it establish that history consistency improves physical control; the static steering experiments below test different operators and outcomes.

4.2 The learned edit improves some forecasts

On recorded development trajectories, the four-direction edit reduces H6 robot-state embedding MSE by 2.360% on Reach and 2.192% on Reach-Wall. These are prediction errors for recorded futures. They do not measure how the robot moves under the edited planner.

An earlier rank-one sweep finds larger H6 forecast corrections at early predictor blocks on both tasks, in FP32 and BF16. Push-T does not show the same pattern. The sweep uses independently fitted rank-one operators; it does not test moving the later four-direction edit to another layer. Appendix B retains the full layer analysis.

4.3 A score-margin bound characterizes fixed-bank selection

We next score identical banks of 300 action sequences in 192 development states, 96 per MetaWorld task. The learned edit changes the winning plan in zero Reach states and one Reach-Wall state.

To assess why a winner can remain stable, let m be the original cost gap between the best and second-best plans. Let δ_i be the cost change for plan i , and let $r = \max_i \delta_i - \min_i \delta_i$ be the range of these changes. If $r < m$, no competitor can overtake the winner: its original disadvantage exceeds any relative improvement the edit can give it. Appendix C gives the proof and all four edit comparisons.

This condition holds for the learned edit in 184 of 192 states: 94 Reach and 90 Reach-Wall. Failure of the bound does not imply a changed winner. Conversely, preserving the best plan does not preserve all ten elites used by CEM. The bound applies to the saved finite set of candidates, and does not certify the result of later search iterations.

A separate 64-state replay asks which part of the edit produces the score changes. Keeping only the correction shared across candidates reconstructs the full edit’s relative cost changes with scores of 0.983 on Reach and 0.998 on Reach-Wall. Random edits show a similar pattern. A shared activation correction can still change relative costs because squared goal distance depends on each candidate’s starting prediction (Equation 7). Components retain their natural magnitudes, so this test does not compare their efficacy at equal energy.

4.4 Iterative search can return different plans

CEM changes its next candidate population using the current elite plans. Small score changes can therefore alter later searches even when the first winner remains unchanged. A 56-state extension, with 28 states per task, runs all fifteen iterations for unsteered, learned, and random-subspace conditions.

Both edits change the returned action prefixes. Mean RMS differences from unsteered are 0.697/0.578 for learned/random on Reach and 0.500/0.217 on Reach-Wall, in normalized action coordinates. All six registered learned-minus-random intervals include zero. These measurements establish changed plans, but do not establish an advantage for the learned edit or measure physical displacement. Complete search curves and contrasts appear in Appendix C.1.


4.5 Physical replay separates forecast accuracy from plan quality

We execute each condition’s selected prefix from the same simulator reset in all 56 extension states. Each trajectory contains fifteen elementary actions. A repeated unsteered execution checks exact reset and trajectory agreement, giving 224 trajectories. All three models also predict all three selected plans, giving 504 model/plan forecasts.

Holding the unsteered-selected actions fixed, the learned edit reduces weighted forecast MSE by 0.002720 on Reach (simultaneous 95% interval $[0.002342, 0.003061]$) and 0.002268 on Reach-Wall ($[0.001905, 0.002607]$). These are absolute error reductions. For executed plans, however, all eight

The first choice stays the same in 191 / 192 states

Learned edit on fixed candidate sets. The curves show all four edits.



Learned edit: 184/192 certified unchanged; 1/192 observed winner changes.

96 contexts/task · same 300 actions per comparison · initial population, not an entire CEM search

Figure 3: Most fixed candidate sets retain their winner. The horizontal axis divides the range of score changes by the original winner’s lead. Values below one certify an unchanged winner. Curves retain all 96 states per task and all four edits; the learned edit changes one winner across both tasks. This comparison holds candidate actions fixed.

175 contrasts in physical end-effector distance and actual encoded-goal cost include zero under the reg-
 176 istered twelve-comparison correction (Figure 4).

177 Thus we can detect better same-action forecasts in this experiment without establishing that the
 178 selected actions bring the robot closer to its goal. This tests short executions on development states.
 179 It is separate from full-task success. CEM optimizes H6 while this follow-up measures the H3
 180 endpoint, so discrepancies here do not by themselves show exploitation of the optimized objective.

181 4.6 Local geometry depends on how the predictor is computed

182 With weights, inputs, and interpolation coefficients fixed across 64 contexts, cubic activation recon-
 183 struction outperforms linear reconstruction at every block in FP32; BF16 reverses the comparison.
 184 Rounding only the FP32 output does not reproduce actual BF16 computation. A separate five-task
 185 interpolation study finds small, mixed downstream forecast benefits after matching requested edit
 186 doses. An apparent geometric advantage therefore needs numerical controls and a separate test of
 187 its effect on prediction (Appendix D).


188 4.7 Protected behavioral evaluation

189 The protected panel evaluates eight arms on 384 fresh scenarios across four tasks, for 3,072 runs.
 190 Table 1 gives every arm. All 48 registered simultaneous intervals include zero, covering compar-
 191 isons with unsteered, learned versus random edits, and component interactions. This is insufficient
 192 evidence for a reliable success gain; it is not an equivalence result.

193 Steering does change which scenarios succeed. On Reach, unsteered and learned steering each
 194 succeed on 52 of 96 scenarios. The edit rescues 21 failures and loses 21 successes. Earlier develop-
 195 ment rates use different scenarios and planner randomness, and are not substitutes for the concurrent
 196 protected baseline.

Better forecasts; no established physical gain

Same actions for prediction tests; each model's chosen actions for execution.



28 paired states per task. Bars: simultaneous 95% intervals across 12 comparisons.

Negative values favor the edit. Short executions on development states.

Figure 4: Forecast accuracy and physical outcome in the same 56 states. Negative effects favor the edit. The top row holds actions fixed and measures prediction error. The lower rows execute each model's selected actions and measure the resulting distance and goal cost. Dots show paired states, diamonds show means, and bars show the registered family-twelve intervals. Learned forecasts improve in both tasks; all eight physical-outcome intervals cross zero.

5 Discussion and limitations

The action-history test identifies a specific source of mismatch between a local patch and an input counterfactual: the same action-derived information is reused at the next prediction step. An evaluation at the edited step alone misses this difference. In longer-horizon agents, repeated observations, cached features, and action histories provide further places to test this issue, but our experiment establishes it only for the two-frame JEPa-WM predictor. H6 is not evidence of long-context scaling, and exact reproduction of a model counterfactual is not evidence of physical correctness.

The planning experiments address a separate requirement: an edit must change candidate-relative costs enough to alter the relevant selection decision. A shared activation correction can do this because its effect on squared goal distance depends on each candidate's original prediction. Candidate-specific edit coefficients are therefore not necessary. The fixed-bank bound explains many unchanged initial winners, while adaptive search measures later plan changes. Neither analysis identifies the cause of protected outcome changes: those episodes used different inputs and randomness and did not retain numerical candidate costs or elite identities.

Our physical follow-up executes three selected plans per state, not all 300 candidates, and measures H3 while CEM optimizes H6. It cannot establish the true physical ranking of the complete candidate bank. The study uses one released checkpoint per task, task-specific fitted edits, and 96

Table 1: Complete protected success rates (%). Each cell has 96 paired scenarios. All 48 registered simultaneous contrast intervals include zero.

Arm	Reach	Reach-Wall	PointMaze	Wall
Native	54.17	29.17	86.46	81.25
Refined rank-four	54.17	20.83	88.54	80.21
Calibrated random subspace	44.79	21.88	86.46	80.21
Equal-budget coupling	47.92	18.75	84.38	83.33
Dose-matched random directions	46.88	27.08	87.50	77.08
Unscaled joint	52.08	22.92	84.38	77.08
Visual only	54.17	22.92	81.25	80.21
Action conditioning only	47.92	17.71	81.25	78.13

protected scenarios per task, without independent training-seed replication. Development diagnostics informed the analysis. All 48 protected contrast intervals include zero, leaving small behavioral effects unresolved. Full tables retain random controls, unfavorable results, and unavailable historical cells.

6 Conclusion

Frozen JEPA-WM lets us test how changes inside a predictor affect the computations a planner uses. Replacing an action-derived condition across both context appearances reproduces a specified input-action counterfactual; a one-time patch agrees initially and diverges later. Learned corrections improve some forecasts, a score-margin bound identifies stable fixed-bank choices, and adaptive search can still produce different plans. Measuring these stages separately gives world-model researchers a way to distinguish a consistent intervention, a better forecast, and a better control decision.

References

- [1] Sonia Joseph, Quentin Garrido, Randall Balestriero, Matthew Kowal, Thomas Fel, Shahab Bakhtiari, Blake Richards, and Mike Rabbat. Interpreting physics in video world models. *arXiv preprint arXiv:2602.07050*, 2026. URL <https://arxiv.org/abs/2602.07050v1>. Version 1, February 4, 2026.
- [2] Miranda Muqing Miao, Subin Kim, Brandon Yang, and Lyle Ungar. Contrastive conceptor activation steering (COAST): Unlocking vision-language-action models through hidden states. *arXiv preprint arXiv:2605.17144*, 2026. URL <https://arxiv.org/abs/2605.17144v1>. Version 1, May 16, 2026.
- [3] Basile Terver, Tsung-Yen Yang, Jean Ponce, Adrien Bardes, and Yann LeCun. What drives success in physical planning with joint-embedding predictive world models? *arXiv preprint arXiv:2512.24497*, 2025. URL <https://arxiv.org/abs/2512.24497v4>. Version 4, September 2, 2026.
- [4] Daniel Wurgaft, Can Rager, Matthew Kowal, Vasudev Shyam, Sheridan Feucht, Usha Bhalla, Tal Haklay, Eric Bigelow, Raphael Sarfati, Thomas McGrath, Owen Lewis, Jack Merullo, Noah Goodman, Thomas Fel, Atticus Geiger, and Ekdeep Singh Lubana. Manifold steering reveals the shared geometry of neural network representation and behavior. *arXiv preprint arXiv:2605.05115*, 2026. URL <https://arxiv.org/abs/2605.05115v1>. Version 1, May 6, 2026.
- [5] Gaoyue Zhou, Hengkai Pan, Yann LeCun, and Lerrel Pinto. DINO-WM: World models on pre-trained visual features enable zero-shot planning. *arXiv preprint arXiv:2411.04983*, 2024. URL <https://arxiv.org/abs/2411.04983>.

Table 2: Eight arms, three alternative intervention sites. Native plus these seven edits form the protected panel. Earlier layer and geometry sweeps are separate. R is never combined with V/A.

Family	Intervention	Comparator / ablation
Refined	Four-direction R at B3 output	Calibrated random subspace
Equal-budget coupling	V+A; each dose scaled by $1/\sqrt{2}$	Same-dose random V+A
Component ablation	Original-dose V+A	V only; A only

A Intervention details and data populations

Frozen predictor and planner. We use released JEPA-WM checkpoints and task-specific fitted interventions. Reach and Reach-Wall share a checkpoint, not their fitted operators. DINO visual features and proprioceptive features feed a six-block action-conditioned predictor. All base encoders and predictor weights remain frozen during intervention; this does not imply that every encoder was frozen during original training.

The analyzed MetaWorld planner uses the actual weighted goal cost

$$C(a_{1:H}) = \text{MSE}(z_H^v, z_g^v) + 0.1 \text{MSE}(z_H^p, z_g^p), \quad (3)$$

with each MSE averaged over its modality’s dimensions. The cost is *not* visual-only. Recorded-future proprioceptive error and goal-directed planning cost nevertheless use different targets, weights, and action inputs. CEM uses 300 candidates. MetaWorld uses 15 optimization iterations and seven planning calls over the recorded 100-elementary-step episodes; navigation uses 30 iterations. Imagined steps, simulator steps, CEM iterations, and replanning calls are distinct units.

Refined rank-four correction. The later fixed-response operator edits the B3 output at H3. With native field h , orthonormal four-direction basis B , and $\phi(h)$ consisting of an intercept and three standardized native-feature scores, an offline map W gives

$$d(h) = BW\phi(h), \quad h' = h + \eta d(h)/\|d(h)\|_2, \quad (4)$$

using the frozen zero-edit rule below the correction threshold. The edit covers all 256 visual patches within the existing batch, without online response probes or native-shadow forecasts. Its random-subspace comparator has separate response calibration and matched support, rank, basis-projector spectrum, and prescribed dose—not identical coefficient-map singular values. This successor is not an equivalent implementation of the earlier online response solver.

Visual/action interventions. A fit-only maximum-covariance direction pair edits the H3 visual predictor input and B3 action-conditioning vector. Unscaled joint, visual-only, and action-only arms form a fixed-dose factorial. Equal-budget coupling scales both standardized doses by $1/\sqrt{2}$; its random-direction comparator uses the same sites and dose policy. This comparison is distinct from the unscaled factorial (Table 2).

Separate populations and units. Fit trajectories supply directions, readouts, and calibration. Separate recorded-action development lineages measure forecast effects and inform choices; fitting-disjoint does not mean selection-protected. Corrected primary offline pools contain 33 Reach, 27 Reach-Wall, and 21 Push-T lineages. Geometry analyses also include 192 Wall and 200 PointMaze lineages. The archived decision diagnostic has 96 canonical development scenarios per MetaWorld task, each with one shared initial bank of 300 action sequences. Protected confirmation instead uses 96 *new* scientific scenarios per task on Reach, Reach-Wall, PointMaze, and Wall. The earlier six-task behavioral table and interrupted native-only cohort remain separate.

The archival CPU analyses recount hash-verified records. New GPU diagnostics separately reexecute the frozen predictor on 64 canonical development contexts (IDs 0–31 per MetaWorld task), with an eight-context actual-CEM pilot (IDs 0–3 per task). Controlled arithmetic uses all 64 contexts; action-conditioning specificity uses 16 (IDs 0–7 per task), each with original and freshly seeded candidate banks. Fresh candidates are not fresh scenarios. These diagnostics do not execute physical actions. Offline effects use source lineages as units; candidate costs and behavior use scenarios. Candidates, heads, horizons, and calls are not independent replicates. Follow-up protocols specify their uncertainty families before outcomes; the studies remain development analyses, not replacements for protected confirmation.




Figure 5: An early-to-late forecast response in the rank-one sweep. H6 proprioceptive-error reduction relative to native in BF16 and FP32, retaining all eight intervention supports and three tasks. The full 576-cell support/modality/horizon grid and matched-random comparisons are provided separately. These are independently fitted rank-one interventions, not a layer sweep of the later rank-four successor.

285 B Layer-by-layer forecast effects

286 A recount verifies all 576 task/precision/support/modality/horizon cells in the earlier layer grid. In
 287 its singleton *rank-one* sweep, H6 proprioceptive-error reduction decreases monotonically from B0
 288 to B5 on both MetaWorld tasks in both numerical conditions (Figure 5). Reach B0/B5 reductions are
 289 3.026%/0.475% in BF16 and 3.050%/0.494% in FP32; Reach-Wall reductions are 1.526%/0.298%
 290 and 1.334%/0.228%. Original BF16 simultaneous B0-minus-B5 intervals are [1.769, 3.333] and
 291 [0.243, 2.214] percentage points of native MSE.

292 This is intervention susceptibility, not a uniquely semantic layer. Layer-specific fits, representation
 293 scale, dose realization, and downstream computation all enter the comparison. Reach B0 versus B1
 294 remains unresolved, and Push-T does not reproduce the MetaWorld H6 benefit. Endpoints before
 295 H3 precede the intervention.

296 The later *rank-four fixed-response* successor is evaluated at B3, not across this layer sweep. It
 297 reduces development H6 proprioceptive MSE by 2.360% on Reach and 2.192% on Reach-Wall
 298 while removing online response probes. Neither result establishes that moving the successor to B0
 299 would improve planning.

Table 3: All diagnostic arms, 96 banks per task. Certified: $r < m$. Changed: actual winning-index changes. RMS is candidate-centered cost change divided by native within-bank cost standard deviation, then averaged over scenarios.

Task	Arm	Certified	Changed	RMS (%)
Reach	Refined	94	0	0.2466
	Random subspace	95	0	0.2116
	Coupling	82	1	0.7625
	Random directions	92	0	0.3461
Reach-Wall	Refined	90	1	0.2786
	Random subspace	89	0	0.2922
	Coupling	81	2	1.0786
	Random directions	88	1	0.4748

300 C Candidate scores, component replay, and search

301 **An exact sufficient certificate.** For a fixed candidate bank, let C_i be native cost, i^* its unique
 302 winner, and $\delta_i = C'_i - C_i$ the edit-induced change. Define

$$m = \min_{j \neq i^*} (C_j - C_{i^*}), \quad r = \max_i \delta_i - \min_i \delta_i. \quad (5)$$

303 If $r < m$, the winner cannot change, because every competitor satisfies

$$C'_j - C'_{i^*} = (C_j - C_{i^*}) + (\delta_j - \delta_{i^*}) \geq m - r > 0. \quad (6)$$

304 This sufficient inequality concerns the saved finite bank, not generalization. Failure to certify does
 305 not imply a flip; actual selection and tie handling are checked separately. The diagnostic selects the
 306 bank’s best candidate. CEM instead updates its proposal from ten elites, so preserving the best index
 307 alone does not certify an unchanged CEM update or final plan.

308 **Measured cost changes are small relative to ranking.** We analyze all 192 completed develop-
 309 ment scenarios and four non-native diagnostic arms (768 paired comparisons). For refined editing,
 310 $r < m$ in 94/96 Reach and 90/96 Reach-Wall banks; winners actually change in 0/96 and 1/96. The
 311 mean scenario-level ratio

$$100 \text{ RMS}_i(\delta_i - \bar{\delta}) / \sqrt{\text{Var}_i(C_i)}$$

312 is 0.2466% and 0.2786%. Coupling certifies 82/96 and 81/96 banks and changes 1/96 and 2/96
 313 winners. Table 3 retains both random controls.

314 **Shared coefficients need not produce shared costs.** A separate decomposition of 768 protected
 315 arm records finds that 98.855% and 99.615% of refined coefficient energy on Reach and Reach-Wall
 316 is shared across candidates within a call. This is coefficient geometry, not the fraction of decision-
 317 irrelevant information. Even a constant final-latent translation d changes squared goal cost by

$$\|z_i + d - g\|^2 - \|z_i - g\|^2 = 2d^\top(z_i - g) + \|d\|^2, \quad (7)$$

318 whose first term varies across candidates. The actual edit also passes through a nonlinear predictor.
 319 Margin measurements therefore add information that coefficient energy alone cannot supply.

320 Archived decision banks and protected episodes differ in inputs and planner RNG. Protected records
 321 did not retain numeric candidate costs or elite identities. Thus the certificate cannot explain which
 322 protected scenarios were rescued or broken. A low first-bank flip count also does not establish
 323 replanning drift as their cause.

324 C.1 Component replay and actual adaptive search

325 **A shared field carries relative-score changes.** We test the shared-correction question on all 64
 326 fixed development contexts. For each learned or calibrated-random correction, cache its native
 327 H3/B3 field and replay the full field, zero field, candidate-common mean, and candidate-centered re-
 328 mainder on the same 300 actions. Cached-full and zero controls reproduce their reference forecasts
 329 and scores exactly. Neither component is renormalized, and no operator is refitted.

Small forecast edits. Different search paths.

Same initial candidate actions. Separate adaptive searches.




Figure 6: Adaptive search can diverge despite nearly shared initial elites. Complete 56-context extension, both edit arms and all fifteen iterations. Shading is one scenario standard error, not a confidence interval. RMS uses all 120 proposal-mean coordinates; the returned prefix uses 60 coordinates. Entropy curves and every paired returned-prefix endpoint are retained in the public CEM report. This is a search measurement, not a physical-quality test.

Let u be the full intervention’s cost change, centered across candidates, and u_k the corresponding change for component k . We report $R_k = 1 - \|u - u_k\|_2^2 / \|u\|_2^2$, retaining undefined zero-denominator cases and allowing negative values. This is a reconstruction score, not a probability. For learned edits, common-only R is 0.9827 on Reach (marginal 95% interval [0.9793, 0.9858]) and 0.9975 on Reach-Wall ([0.9971, 0.9980]); centered-only values are 0.0218 and 0.0135. Common-field energy fractions are 99.53% and 99.78%. Random-subspace common-only scores are similarly high, 0.9795 and 0.9854. These results establish the common component’s effect at its natural energy, not superiority over equal-energy centered steering.

Small initial changes can accompany divergent adaptive searches. After an initial eight-case pilot, a separately frozen extension reruns all three arms on the same receiving GPU for 56 new development contexts ($n = 28$ per task). Each search uses fifteen CEM iterations, 300 candidates and ten elites; traced/untraced selected-plan and RNG parity pass for all 336 searches. The original eight cases remain separate from primary inference. Every arm starts from identical candidate actions, but initial elite sets need not be identical: mean learned/random overlap with the native top ten is 9.893/9.857 on Reach and 9.964/9.929 on Reach-Wall.

Later proposals adapt to each arm’s scores, so later candidate IDs are not matched actions. Figure 6 retains all fifteen iterations. Returned three-step, 60-coordinate prefixes have mean RMS differences from unsteered of 0.697/0.578 for learned/random edits on Reach and 0.500/0.217 on Reach-Wall. These normalized model-action differences are not physical displacements or success gains. All six prespecified learned-minus-random intervals include zero (Table 4); the positive marginal Reach-Wall prefix contrast does not survive the registered family correction. Search-path changes therefore extend beyond the original eight cases, without establishing greater divergence or a consistent entropy effect for learned edits over calibrated random edits.

C.2 Forecast improvement and physical prefix quality remain distinct

The separately frozen physical follow-up covers all 56 extension contexts, $n = 28$ per task. Each unsteered, learned, and calibrated-random selected prefix executes from the original seeded simulator

Table 4: All six registered adaptive-search contrasts. Learned minus calibrated random, relative to concurrent unsteered; $n = 28$ paired contexts per task. Intervals use 20,000 paired-context bootstrap draws and Bonferroni family six. Proposal and entropy effects average all fifteen iterations.

Task	Metric	Mean	Simultaneous 95% interval
Reach	Prefix RMS	0.1197	$[-0.1284, 0.3487]$
	Proposal RMS	0.0403	$[-0.0717, 0.1353]$
	Entropy (nats)	0.2631	$[-0.8152, 1.4570]$
Reach-Wall	Prefix RMS	0.2827	$[-0.0070, 0.5294]$
	Proposal RMS	0.0809	$[-0.0541, 0.1929]$
	Entropy (nats)	0.3596	$[-0.6327, 1.3388]$

reset for fifteen elementary actions. A repeated unsteered trajectory checks byte-identical observations and physics, giving 224 trajectories in total. Every model forecasts every plan, retaining 504 scientific model/plan evaluations. Exact reset, native-repeat and H3 future-tail receiving checks pass; the preprocessing amendment and failed first attempt remain preserved in the accompanying physical-prefix report.

On the same unsteered-selected actions, learned-minus-unsteered weighted forecast MSE is -0.002720 on Reach (simultaneous 95% interval $[-0.003061, -0.002342]$) and -0.002268 on Reach-Wall ($[-0.002607, -0.001905]$). These are absolute errors, not percentages. All eight physical end-effector-distance and encoded-goal-cost contrasts for the two edit arms include zero under the prespecified twelve-contrast family correction (Figure 4). Thus this experiment establishes improved same-action forecasts in both tasks without an established physical-prefix gain. CEM optimized H6, whereas execution and scoring here cover H3; predicted/actual H3 preference disagreement alone does not establish exploitation of the H6 objective. This is development-prefix execution, not full-task success or another protected evaluation.

C.3 Action conditions affect rankings and elite membership differently

In a separate 16-context experiment, each candidate receives the next candidate’s newest action-derived AdaLN condition at H3, one block at a time. A private-generator isotropic perturbation matches each donor-minus-recipient L2 norm. All six blocks and both original and freshly seeded 300-action banks are retained. This is donor interchange within a context, not an intervention on explicit action tokens or semantic physical labels.

At B1 and B4, donor substitution disrupts native official-cost rankings more than the random edit: excess Spearman loss is 0.0071 – 0.0084 across the two tasks and two banks. Six-layer simultaneous intervals exclude zero in these cells, but no layer reaches the prespecified 0.05 practical flag. Native rank agreement remains at least 0.98979 in every official-cost cell. At B1, however, donor edits lose a mean 1.0 – 1.125 additional members of the native top ten relative to random edits. Small global rank changes and a changed elite set can coexist; neither establishes physical action quality.

Delivered-norm mismatch is not supported here: the maximum relative error is 2.72×10^{-8} . Distributional mismatch remains. The affine action encoder maps 20 input coordinates into 400 condition coordinates; isotropic noise need not lie in its at-most-20-dimensional reachable subspace. Matching total L2 therefore does not match action-reachable dose. Complete specificity curves and uncertainty appear in Appendix F; the following separately frozen study adds range controls.

C.4 Coherent action changes require consistent history

An H3 action appears twice in the predictor’s two-frame context: newest at H3, then older at H4. On the same 16 contexts and both fixed banks, we compare coherent raw H3 action substitution with donor-condition replacement at H3 only or at both appearances. All-block persistent replacement reproduces the raw-action counterfactual byte-exactly through H6. All-block H3-only replacement matches at H3 but has H6 reconstruction $R = 0.5029$ – 0.5079 across task/bank cells (Figure 8). Here $R = 1 - \text{MSE}(\hat{z}_{\text{patch}}, \hat{z}_{\text{cf}}) / \text{MSE}(\hat{z}_{\text{native}}, \hat{z}_{\text{cf}})$, using ratios of candidate means under the official modality weighting. This is a predictor-internal counterfactual, not its physical correctness.

Predicted endpoint versus the endpoint reached

Same three plans, evaluated by each model and executed after identical resets.





Figure 7: The complete prediction/execution grid. All three models predict all three selected plans in every context: 504 points. Axes use visual goal MSE plus 0.1 times proprioceptive goal MSE, with equal scales within task. Dashed lines mark equal predicted and actual encoded H3 goal cost. These are endpoint goal costs, distinct from prediction-versus-endpoint forecast error.

395 The six-layer contrasts retain a 72-cell simultaneous primary family (two tasks, two banks, six layers,
 396 three contrasts). Persistent versus H3-only reconstruction is positively separated at B0–B3 in both
 397 tasks/banks, 16 of 24 cells; B4/B5 remain unresolved. The largest observed persistence difference
 398 occurs at B1, $\Delta R = 0.4991$ – 0.5057 . H3 donor replacement exceeds equal-norm range-random
 399 reconstruction at B0/B1/B2/B4, but not established at B3/B5. These are a complete depth map, not
 400 a retrospectively selected intervention site.

401 Range-random edits disrupt native rankings more than off-range random edits at B1/B4 in both
 402 banks, by about 0.0021–0.0025 Spearman units; only fresh-bank B2 adds positive separation. The
 403 frozen encoder has rank 20 of 400; isotropic controls place 4.995% mean energy in its range. Actual
 404 norm error is at most 3.41×10^{-8} . Input-range membership is not a global activation manifold,
 405 and off-range does not mean causally inert. Together, these interventions identify an action-history
 406 dependency within this predictor. They neither reproduce the older static steering operators nor
 407 establish the cause of protected outcome changes.

An action returns in the next context window.

A one-time patch is not the same as changing that action.




All six blocks patched. The donor action reappears as history at H4.

Solid / dashed: original / fresh action banks · all 16 contexts, $n=8/\text{task}$ · full marginal 95% bands

Predictor-internal action counterfactuals, not physical outcomes or protected-confirmation mechanisms.

Figure 8: A one-time internal patch is not a coherent action change. The donor H3 condition must also replace the older occurrence at H4. All six blocks, both original/fresh action banks, and all 16 contexts are retained. The exact persistent/raw-action match is a computational consistency control; the H3-only deviation is measured. Bands are full marginal context-bootstrap intervals. No physical actions are executed.



All 72 registered contrasts · simultaneous 95% intervals
 n = 8 paired development contexts/task · no physical outcome

Figure 9: All 72 registered action-counterfactual contrasts. All six predictor blocks, both tasks, and both original/fresh banks are retained. The top comparison measures persistent minus H3-only reconstruction. The other comparisons retain the action-encoder range controls. Intervals use the registered simultaneous correction; these differ from the marginal bands in Figure 2.

C.5 Action-history consistency and candidate selection

The ranking reanalysis uses the original H6 goal-cost arrays, $\text{MSE}_v + 0.1 \text{MSE}_p$ in their archived FP32 arithmetic, with 300 shared candidate IDs per bank. The primary reference is the coherent raw H3 action substitution. At each of the six single-layer sites and the all-block site, we compare H3-only and persistent patches with that reference. Secondary comparisons retain persistent versus H3-only directly and native versus the coherent reference. The new analysis protocol was fixed before calculating these pairwise metrics; the original reconstruction outcomes were already known, so this is post hoc development analysis.

Spearman correlation uses average ranks for exact ties. Elite overlap is the intersection of the actual archived top-ten sets, whose validity is checked against their costs. Winner comparisons use the lowest candidate ID among equal minimum costs, and we separately audit minimum-set overlap and top-ten cutoff ties. Every minimum is unique and no cutoff is tied in these records. There are no undefined rank correlations. Only H6 cost arrays were preserved; the H3/H4 reconstruction endpoints cannot supply intermediate rankings.

Table 5: All-block H3-only agreement with the coherent raw-action reference. Each row contains eight contexts. Persistent all-block patches have Spearman 1, overlap 10, and zero different winners in every row. Original/fresh banks share contexts.


Task	Bank	Spearman	Shared elites / 10	Different winner / 8
Reach	Original	0.8775	6.625	3
Reach	Fresh	0.8814	6.875	5
Reach-Wall	Original	0.8940	5.750	5
Reach-Wall	Fresh	0.8907	6.875	4

Persistent all-block equality is the prior computation-parity control. The additional observation is that omitting the historical occurrence changes not only forecast reconstruction but also rankings, elite membership, and the winning candidate. Figure 10 retains all sites. B1 has a persistent-minus-H3-only Spearman difference of 0.1249–0.1371 and 1.75–3.125 additional shared elites per context. Its net gain in reference-winner matches is 0/8 and 3/8 on Reach (original/fresh) and 1/8 and 2/8 on Reach-Wall. Direct patch-to-patch winner changes are 3/8, 4/8, 2/8, and 4/8 respectively: a changed winner need not move toward the reference. Mean rank agreement improves at B0–B4 in every bank; B5 has small negative differences, from -0.000154 to -0.000083 .

All 704 pairwise records, 224 paired-difference records, and 84 descriptive mean differences are retained. Contexts are the units; banks and layers are shown separately. No new significance tests, confidence intervals, candidate filtering, or layer selection are performed. An independent SciPy/NumPy calculation agrees with all pairwise metrics. The input export binds 512 cost vectors to 80 hash-verified original archive members and sixteen cloud readback receipts. Agreement with a model counterfactual does not establish better physical plans.

Action history changes which candidates the model prefers.

Agreement with the coherent input-action change at H6 · all blocks and all six single-layer sites



Original: solid / square · fresh: dashed / circle · eight contexts per task · descriptive means

All-block persistent matching is the exact control. Single-block persistence need not preserve the same winner.

Figure 10: Action history changes which candidates the model prefers. Mean agreement with the coherent raw-action reference at H6, retaining both tasks, both original/fresh banks, and all seven intervention sites. Each mean contains eight contexts. Lines show descriptive values without new inferential intervals. Persistent all-block matching is exact; single-layer persistence has mixed effects on winner agreement.

Precision changes what looks locally curved

Same weights, same action perturbations, every predictor block.




Figure 11: The local reconstruction comparison depends on computation, not only output rounding. Same weights, inputs, action perturbations and coefficients; 32 contexts per task, all six blocks. Curves show geometric mean cubic/linear activation-MSE ratios with full marginal 95% bands. The prespecified paired arithmetic contrasts use six-layer correction. These are reconstruction errors, not forecast or robot-success gains.

D Numerical precision and activation reconstruction

A controlled arithmetic comparison. The 64-context experiment fixes model weights, encoded inputs, the literal first action coordinate, four symmetric perturbations ($\pm 0.1, \pm 0.05$), and reconstruction coefficients. At every block, it compares strict FP32, actual BF16-autocast computation, and FP32 fields rounded to BF16 then restored to FP32. Each condition reconstructs its own unperturbed center; no fitted bank or dose-normalized correction changes between conditions.

FP32 favors cubic reconstruction at all six blocks: mean log cubic/linear MSE ratios span -5.40 to -7.89 . Output round-tripping moves these ratios near equality (-0.098 to $+0.010$), whereas actual BF16 produces $+0.160$ to $+0.306$. The actual-BF16-minus-rounded contrast is $+0.224$ to $+0.336$ across all twelve task/block cells, with six-layer-corrected lower bounds above the prespecified 0.1 log-unit margin. Thus output rounding degrades the apparent reconstruction advantage but does not quantitatively reproduce the actual BF16 result. The experiment isolates numerical-condition sensitivity, not an individual kernel or a physical manifold.

Precision-sensitive local interpolation. The action-response experiment withholds the center of a four-donor activation path. Coordinates $-1, -\frac{1}{2}, \frac{1}{2}, 1$ index frozen local action offsets. Equal-anchor linear reconstruction averages the donors. The archived cubic estimator at zero is

$$\hat{h}_{\text{center}} = -\frac{1}{6}h_{-1} + \frac{2}{3}h_{-1/2} + \frac{2}{3}h_{1/2} - \frac{1}{6}h_1. \quad (8)$$

At this symmetric evaluation point the same functional is obtained by quadratic least squares. It does not uniquely identify third-order dynamics or a dense activation manifold.

Across all five tasks, raw omitted-activation MSE is 35.15 – 46.77% *higher* than linear reconstruction in BF16, but 93.96 – 99.99% *lower* in FP32. This does not correspond to a similar improvement in future prediction. Under matched requested doses, the H6 proprioceptive-MSE advantage over the linear arm is small and mixed: BF16 task means range from -0.0213% to $+0.0264\%$ of native future MSE; FP32 ranges from -0.0141% to $+0.0051\%$. The two endpoints have different targets and denominators (Appendix H), so these scales are not an attenuation coefficient.

460 **Requested dose is not exact realized-energy equivalence.** The precision conditions also differ
 461 in fitted banks and prescribed dose, preventing attribution solely to rounding. In FP32 Reach/Reach-
 462 Wall, the cubic arm’s per-window realized-energy tolerance passes only 49.0%/50.7% of windows,
 463 versus 98.1%/98.4% for linear reconstruction, despite small mean realized/requested L2 discrepan-
 464 cies. The comparison is therefore labelled by its actual requested-dose contract, not perfect per-
 465 window energy matching.

466 **Output nonadditivity differs from loss interaction.** For the unscaled visual/action factorial, let
 467 $e = z_0 - y$, $v = z_V - z_0$, $a = z_A - z_0$, and $q = z_{VA} - z_V - z_A + z_0$. Squared-error interaction
 468 obeys

$$L(VA) - L(V) - L(A) + L(0) = 2v^\top a + 2(e + v + a)^\top q + \|q\|^2, \quad (9)$$

469 with the endpoint’s dimensional averaging. Thus loss interaction can be nonzero even when output
 470 interaction q vanishes.

471 Saved H6 output-interaction squared norms, divided by native proprioceptive MSE, are 0.419%,
 472 0.512%, and 5.020% in BF16 for Reach, Reach-Wall, and Push-T. Corresponding FP32 values are
 473 $4.82 \times 10^{-6}\%$, $1.07 \times 10^{-5}\%$, and 0.00256%. These establish precision-conditioned differences in
 474 retained scalar diagnostics, not that rounding caused them. Missing horizons remain missing; origi-
 475 nal tensors are not reconstructed from scalars. Together, these results distinguish local interpolation
 476 quality, the delivered activation intervention, and its downstream forecast utility. The loss terms can
 477 nearly cancel: BF16 Push-T has a +2.460% additive cross-term and −2.509% nonadditive remain-
 478 der, yielding only −0.049% net H6 loss interaction. Figure 14 retains all seven source conditions,
 479 not just this illustrative case.

480 E Native attention and diagnostic evidence




Figure 12: Native H6 attention distance, all six blocks and sixteen heads. Each cell averages 32 devel-
 opment contexts per task with the same archived zero-action candidate. Colors are shared across tasks. The
 complete H1–H6 grid is retained in the accompanying data. Attention distance is not causal importance or
 evidence of a planning-emergence zone.

481 The observer measures post-RoPE attention-weighted spatial and temporal distance while retaining
482 the native attention kernel and exact output/RNG parity. It covers six blocks, sixteen heads and six
483 horizons on all 64 contexts. Actions enter through AdaLN conditioning rather than extra attention
484 tokens; zero action-token mass is not evidence that actions are ignored. The zero-action input does
485 not characterize attention over arbitrary candidate plans.


486 Native and steered CEM traces retain actual proposal standard deviations. Differential entropy is
487 $\sum_{j=1}^{120} (\log \sigma_j + \frac{1}{2} \log(2\pi e))$, before clipping and insertion of the mean candidate. It depends on
488 action units and may be negative. It is neither attention entropy nor calibrated uncertainty about
489 physical outcomes.

490 All files emitted by these completed diagnostics are cloud-verified. Full H6 tensors were retained for
491 the first eight contexts; later contexts retain compact diagnostics after on-GPU parity checks, under
492 the fixed retention policy. Preservation does not recover tensors that were never saved. The eight
493 actual-CEM follow-up records are fully preserved. An engineering-only amendment changed an
494 incorrect fifteen-callback guard to thirty: each iteration forecasts both the 300-candidate population
495 and its updated mean. Failed attempts remain excluded, with unchanged seeds and scientific arms.

496 **F Action-conditioning specificity**

497 The subsequent action-counterfactual experiment retains all 72 primary contrasts in the accompa-
498 nying `action_counterfactual_primary_contrasts.csv`. Its 20,000 paired bootstrap draws
499 share context weights across banks and layers. Each task/contrast family uses the maximum abso-
500 lute centered deviation over twelve bank/layer cells, with quantile $1 - 0.05/6$ across six families.
501 Undefined cells would remain undefined; all registered cells were defined. All 1,120 forecasts,
502 source identities, cloud receipts, norm controls and full-horizon parity checks passed. The full-layer
503 figure and reproduction receipt are provided alongside the focused timing figure.

Action alignment versus norm-matched random edits



Positive: permutation disrupts native rankings more than random.

All 16 cases; $n=8/\text{task}$. Six-layer simultaneous 95% bands.

Official visual + 0.1 proprio goal cost. No actions executed.

Figure 13: Small global ranking changes can alter the elite set. Complete 16-context donor-interchange experiment, eight contexts per task, two fixed candidate banks, and all six layers. The contrast is excess loss of native goal-cost rank agreement relative to an equal-norm isotropic edit. Six-layer simultaneous intervals are within task/bank/metric families, not across every secondary endpoint. No layer reaches the prespecified 0.05 excess-Spearman practical flag. Rankings are model predictions, not physical counterfactual values.

504 Zero-capture passes reproduce H1–H6 forecasts and costs exactly; interventions leave the earlier
 505 condition time unchanged. All 115,200 delivered norms pass. The 448 total forecasts include
 506 32 engineering passes and do not create 448 independent observations. An import-only receiv-
 507 ing amendment initializes the native planner dependency before the RNG boundary; the final guard
 508 still requires unchanged global RNG throughout all scientific forecasts. The failed attempt remains
 509 archived. The public source-bound receipt exports all 1,152 case metric rows and 144 contrast rows.

Table 6: Historical development only, not pooled into protected confirmation. DROID is a recorded-plan action score over 64 endpoints, not 96 physical robot episodes. MetaWorld component rows marked * use concurrent native 45.83/29.17, not the canonical native row.

Arm	Reach	R.-Wall	Push-T	P.Maze	Wall	DROID
Native	44.79	30.21	59.38	80.21	76.04	51.10
Refined	51.04	31.25	59.38	85.42	78.13	51.05
Random subspace	50.00	23.96	61.46	79.17	76.04	51.00
Coupling	48.96	26.04	61.46	77.08	82.29	51.07
Random directions	60.42	28.13	60.42	84.38	77.08	51.27
Joint	52.08*	29.17*	61.46	79.17	78.13	50.85
Visual	50.00*	38.54*	60.42	81.25	73.96	50.83
Action	42.71*	36.46*	58.33	86.46	76.04	51.11

Table 7: Author-reported benchmark context [3]. Tables 2/11 report late-epoch recipe aggregates; Table 12 reports final checkpoints. Different aggregation, checkpoints and scenarios prevent pairing these rows with our interventions.

Published model	Reach	R.-Wall	Push-T	P.Maze	Wall	DROID
DINO-WM	44.8	35.1	66.0	81.6	64.1	39.4
JEPA-WM recipe, L2	58.2	41.6	70.2	83.9	78.8	48.2
JEPA-WM recipe, L1	55.1	40.8	63.4	79.7	46.7	47.2
JEPA-WM final, L2	49.0	29.2	69.4	83.3	80.9	46.5


511 H Local geometry: targets and missingness

Table 8: Reconstruction reduction is $100(L_{\text{linear}} - L_{\text{cubic}})/L_{\text{linear}}$ for omitted-activation MSE. H6 advantage uses the future proprioceptive loss difference divided by *native future MSE*. Positive favors cubic. These denominators are not interchangeable; values are descriptive means.

Task	<i>n</i>	Reconstruction reduction (%)		H6 advantage (%)	
		BF16	FP32	BF16	FP32
Reach	33	-46.4289	94.1933	0.009638	-0.000231
Reach-Wall	27	-46.7716	93.9634	0.026390	0.000000123
Push-T	21	-35.1464	99.9161	-0.021238	0.005017
Wall	192	-43.3160	99.9748	-0.006288	-0.000688
PointMaze	200	-42.2175	99.9858	0.015918	-0.014072

512 All ten task/precision geometry sources are retained. The seven coupling reports are not a complete
 513 five-task-by-two-precision factorial: registered PointMaze coupling and FP32 Wall coupling are
 514 absent. Output decomposition is unavailable at H2/H4/H5. The compact Wall source lacks per-
 515 arm lineage losses, so its aggregate loss interaction is retained without invented lineage contrasts.
 516 Saved interaction norms and cross terms are recounted and identity-checked, not reconstructed by
 517 reexecuting output tensors. Precision-specific fits and realized doses prevent isolating rounding as a
 518 treatment.

Same-input forecast interaction: D = cross + remainder



Different y-scales. Remainder already includes output-interaction q^2 .

Wall D: aggregate only, no CI. Other intervals: marginal paired 95%.

No PointMaze or FP32 Wall source. This is not a success factorial.

Precision-specific fits and doses preclude a precision-causal claim.

Figure 14: Output pathways and squared-loss interaction are distinct. All seven registered H6 proprioceptive conditions. The additive cross-term plus nonadditive remainder equals the loss factorial; the remainder already includes $\|q\|^2$, which is not added twice. Wall's aggregate-only factorial has no invented interval; its remaining intervals retain the saved fixed native denominator. Separate axis scales, precision-specific fits, and doses preclude treating the panels as a controlled precision experiment. These are fixed-input forecasts, not protected behavioral interactions.

I Reproducibility and analysis boundaries

The protected output is `reports/fresh-confirmation/report.json`, SHA256 8123d71497835fc164f09f5094c4303086 47a3ee2462c66746baea32633a0b15. An independent audit checks 32 rates, 384 unique scenario owners, 3,072 arm hashes, and referenced engineering bundles. The registered bootstrap seed is 2026091221. Same-device pairing and scenario-specific RNG hold within each protected scenario, not between old and fresh cohorts. Input-overlap protection is relative to audited project exposure records, not all base-model pretraining.

The decision reanalysis is `paper/data/decision_geometry.json`, its per-scenario CSV, and `analysis/mechanism/decision_geometry.py`. Every source record is hash-bound. Exploratory marginal scenario-bootstrap intervals use 20,000 resamples and seed 2026091321; the finite-bank certificate itself is arithmetic, not statistical.

Layer measurements are in `paper/data/layer_mechanism_grid.csv`. The pathway recount is `paper/data/pathway_geometry.json`, with lineage, summary, and cross-precision exports. Curated baseline provenance and figure captions are in `benchmark_comparison_sources.json` and `benchmark_figure_provenance.json` under `paper/data`. Protected rates are never recovered from pixels or historical denominators.

Execution and scaling. The frozen PyTorch inference workloads used independent whole-scenario GPU replicas, with 300 candidates batched inside each forward call. All eight arms of a scenario stayed on the same physical device. Pinned cloud staging and asynchronous hash-verified preservation overlapped independent execution; they did not alter scientific inputs. Strict FP32 and TF32-off settings were retained. Evaluation uses independent whole-scenario replicas; the model fits on each device and requires no backward pass. Changes to precision, kernels, or parallelism require replay parity checks before scientific comparison.

Identifying repository and storage links are omitted for anonymous review. The aggregate tables and source-bound figure scripts form the reproducibility package; full raw archives are not bundled in this PDF. New replay or adaptation work requires separately specified inputs, parity checks, and an evidence label. Unfinished experiments are not counted as results here.

Broader impact. The evaluation separates prediction quality, counterfactual consistency, and control outcomes, which can help researchers diagnose a world model before drawing conclusions about its use in an agent. Treating better forecast error as evidence of effective or safe physical control could encourage premature deployment; these simulator results do not support that inference. No new real-robot deployment or human-participant data collection is reported.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction distinguish counterfactual reconstruction, forecast error, plan selection, and physical outcomes. They explicitly limit the context experiment to a two-frame predictor and do not claim context-length scaling.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 5 discusses the context window, task and checkpoint scope, statistical units, missing protected candidate costs, and the absence of training-seed replication.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Appendix C states the fixed-bank and unique-winner assumptions and proves the sufficient score-margin certificate. The goal-distance translation identity is derived there as well.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [No]

Justification: The paper specifies the interventions, populations, and analysis, but exact end-to-end replay also depends on pinned runtime, checkpoint, and input archives not included in this anonymous PDF. Appendix A and the reproducibility appendix describe this boundary.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility.

In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [No]

Justification: The PDF does not supply an anonymous open-access code and raw-data bundle. It retains full aggregate tables and source hashes; identifying repository and private storage links are omitted for review.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Section 3 and Appendix A describe the frozen model, intervention sites, fitted operators, planning settings, distinct fit/development/protected populations, and task-specific evaluation endpoints.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Section 3, the experiment appendices, and figure captions specify context or trajectory resampling, paired comparisons, bootstrap draws, and simultaneous correction. The new history curves retain marginal intervals, explicitly distinguished from simultaneous layer contrasts.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [No]

Justification: The paper reports workload sizes and numerical execution conditions but does not yet provide a complete per-experiment accounting of GPU types, memory, elapsed time, and total failed or preliminary compute.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [N/A]

Justification: The final author attestation to the NeurIPS Code of Ethics has not been completed for this draft.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The reproducibility appendix discusses the potential value of more informative world-model evaluation and the risk of treating improved forecasts as evidence of safe or effective physical control.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: This paper does not release a new general-purpose model or a scraped dataset with elevated misuse risk. It reports bounded simulator and recorded-data analyses of existing checkpoints.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [No]

Justification: The principal world-model and steering sources are cited, but the manuscript does not yet enumerate the license and terms of every third-party code, model, and dataset asset. Raw assets are not redistributed in this PDF.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [No]

Justification: Aggregate analyses and figure scripts are documented in the working repository, but a separately anonymized asset package has not been included with this draft. The manuscript describes its evidence and reproduction boundaries.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: No new crowdsourcing or human-subject experiment was conducted. The physical evaluations use simulators; DROID contributes only an existing recorded-action endpoint.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The study does not recruit human participants or conduct a new human-subject experiment.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: An LLM is not a component of the tested predictor, steering operators, planner, or evaluation endpoints. Assistance with code and manuscript preparation is not presented as a scientific method contribution.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.